

The Full Map: Tokenized Whisky Cask Infrastructure

0. EXPLICIT RULES (LOCKED FOR DESIGN PHASE)

These rules are the canonical constraints for every design doc in this repository:

1. Oracle scope is **physical attributes + reserve status + lifecycle provenance**. This is **not** a NAV oracle.
2. Dollar valuation is optional Tier 3 reference data and is not required for core onchain infrastructure.
3. API runtime is **TypeScript + Hono on bun** for hackathon scope. No mixed FastAPI plan in active docs.
4. Proof of reserve uses `TOKENS_PER_CASK = 1000` vault units to normalize cask count vs ERC-1155 token units.
5. Confidential mode must not reveal raw inventory counts onchain. Only `isFullyReserved` + attestation hash.
6. Baseline simulation mode may include raw counts for debugging, but privacy track narrative must use confidential mode semantics.
7. `FULL_MAP.md` is the source of truth. Other docs must mirror terminology, workflow names, and assumptions.

1. THE REALITY OF CASK DATA

What data actually exists for a whisky cask?

| CASK LIFECYCLE | DATA GENERATED | SOURCE | CONFIDENCE |
|-------------------------------|--|-------------------------------------|--|
| Distillation | New make spirit gauge (volume, proof, date) | DSP records (TTB) | Hard fact |
| Filling | Entry gauge: volume, proof, cask serial, cask type, fill date, warehouse ID | Warehouse records (TTB-required) | Hard fact |
| Years 0-6 (maturation) | NOTHING NEW. Cask sits. No measurement. Angel's share is happening but nobody's measuring it. | - | - |
| Year 6+ (every 2-3yr) | Regauge (if owner requests): RLA, ABV, bulk litres, date, method | Warehouse (on request) | Imprecise Wet dip = in a barrel ±5-10% var Full disgorgement accurate but |
| Transfer | Re-gauge at origin + destination New warehouse ID, date | Both warehouses (TTB-required) | Hard fact (gauges in) |
| Bottling decision | Cask selected, final gauge | DSP/warehouse | Hard fact |
| Bottling | Cask emptied, final yield, bottles produced, proof | DSP records (TTB-required) | Hard fact |
| Sale/Transfer of ownership | Ownership transfer record (warehouse receipt assignment) | Legal docs | Hard fact |

Key insight: There are massive data gaps during maturation.

Between fill and first regauge (potentially 6+ years), the ONLY data is:

- Entry gauge (one measurement from years ago)
- Mathematical estimate of angel's share (model, not measurement)
- The cask physically exists (warehouse can confirm)

Between regauges (2-3 year intervals after year 6):

- Last regauge data (stale within months)
- Angel's share estimate since last regauge

Regauge accuracy is honestly poor for wet dip:

- Cask sizes vary ±10% from nominal (a "250L hogshead" could be 240-270L)
- Temperature affects density readings
- Operator technique matters
- ABV measurement has decimal-place variability
- Only full disgorgement (emptying the cask) is truly accurate, and that's rare

2. THREE DATA TIERS

Tier 1: Verifiable Facts

Things the warehouse/TTB records prove. Ground truth.

| Data Point | Source | Update Frequency | Accuracy |
|-----------------------------|---------------------|--------------------------|---------------------------------------|
| Cask exists | Warehouse inventory | Continuous | Binary — yes/no |
| Cask type | Entry records | Once (immutable) | Exact |
| Fill date | Entry records | Once (immutable) | Exact |
| Entry gauge (volume, proof) | TTB-required gauge | Once | High |
| Location (warehouse ID) | Warehouse system | On transfer | Exact |
| Ownership | Warehouse receipt | On transfer | Exact (legal doc) |
| Last regauge data | Regauge report | Every 2-3yr (after yr 6) | Moderate (wet dip) to High (disgorge) |
| Lifecycle state | Warehouse records | On event | Exact |
| Bottling yield | DSP records | Once (at bottling) | Exact |

Tier 2: Computed Estimates

Math applied to Tier 1 data. Transparent methodology, clearly labeled.

| Data Point | Inputs | Method | Accuracy |
|-------------------------|--|-------------------------|--------------------------|
| Current volume estimate | Last gauge + time elapsed + climate rate | Exponential decay model | ±10-15% between regauges |
| Current age | Fill date + now | Subtraction | Exact |
| Estimated bottle yield | Volume estimate + cask type + quality factor | Linear model | ±15-20% |
| Angel's share to date | Entry gauge + age + climate rate | Compound annual loss | ±5% cumulative |

Tier 3: Market Opinions

Models, comps, vibes. Useful as reference, not truth.

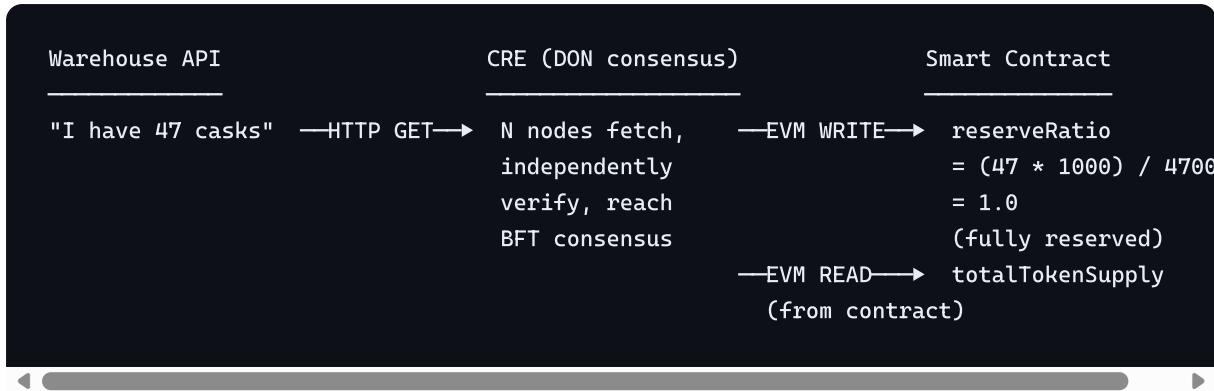
| Data Point | Source | Reliability |
|------------------------------|---------------------------------------|------------------------------------|
| Comparable auction prices | Auction houses (sparse, inconsistent) | Low-moderate (thin market) |
| Whisky index values | RW101, Knight Frank (quarterly) | Moderate (aggregate, not per-cask) |
| Age-based appreciation curve | Our model (calibrated to indices) | It's a model. Labeled as such. |
| Cask type premium | Market convention + auction data | Directionally correct, imprecise |
| "What is this cask worth" | All of the above blended | An informed opinion, not a price |

3. WHAT WE ACTUALLY BUILD (GROUNDED)

Component 1: Proof of Reserve (Tier 1 only — strongest component)

What it proves: Physical inventory backing is at least the outstanding tokenized claims, using a fixed conversion of `TOKENS_PER_CASK = 1000`.

Data flow:



Why CRE matters here (specifically): - Without CRE: Tim's server says "47 casks." Tim also issued the tokens. Self-attestation. - With CRE: N independent nodes fetch the warehouse data, independently verify, reach consensus. The attestation is "the DON confirms the warehouse reports 47 casks." Still trusts the warehouse (which is TTB-regulated), but the PIPELINE is decentralized. - With Confidential HTTP: Same attestation, but the actual number "47" never appears onchain or in node logs. Only "casks ≥ tokens" boolean. Competitors can't extract inventory intelligence.

Frequency: Hourly (or configurable). This is checking inventory count, not measuring liquid.

Attestation modes (explicit): - Baseline simulation mode: write count-derived ratio fields for operator debugging. - Confidential mode (privacy track): write only `isFullyReserved`, `timestamp`, and `attestationHash`.

Honest limitation: This proves cask COUNT matches tokens. It does NOT prove what's IN the casks. A warehouse could have 47 empty barrels. The trust model is: TTB-regulated warehouse + CRE-verified data pipeline. Not trustless end-to-end.

Component 2: Physical Attribute Oracle (Tier 1 + Tier 2)

This is intentionally not a NAV oracle.

What it puts onchain:

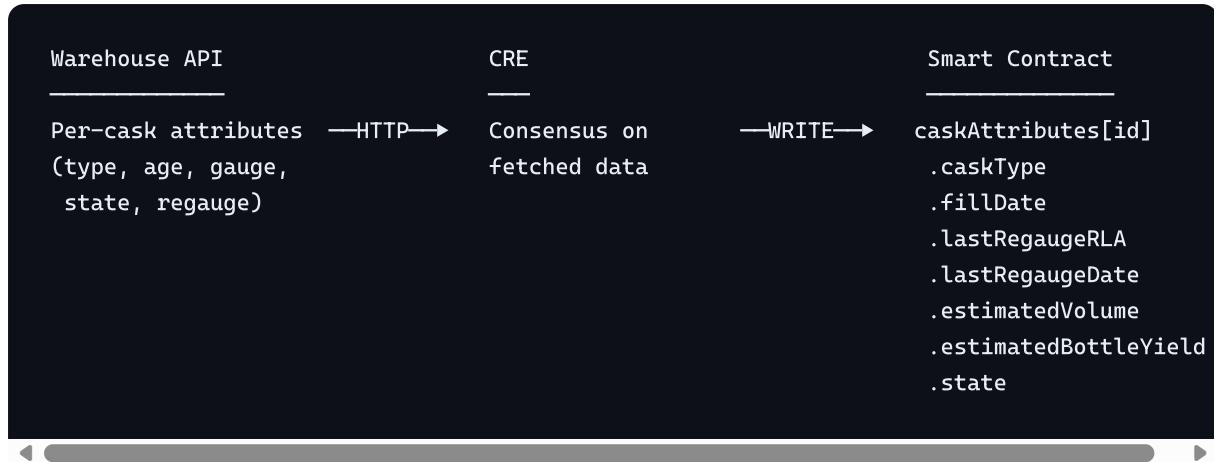
Per cask (from Tier 1 — warehouse records): - Cask type, fill date, entry gauge, last regauge data + date, warehouse ID, lifecycle state

Computed (Tier 2 — transparent math): - Current age (trivial) - Estimated current volume (angel's share model since last gauge, with model parameters visible) - Estimated bottle yield

Aggregate: - Total casks, total estimated volume, portfolio age distribution

What it does NOT put onchain: - A dollar value. Not our job to say what a cask is "worth." - Or if it does, it's CLEARLY labeled as "reference valuation model v1, methodology X, not a market price."

Data flow:



Why this matters: - Any protocol or investor can read verified physical attributes - A lending protocol uses these to set their OWN collateral haircut - A secondary market prices tokens based on visible attributes - Nobody is forced to trust OUR valuation model — they have the data to build their own

Frequency: Daily or on-regauge-event. Attributes don't change fast (age increments daily, volume estimate drifts slowly).

Component 3: Lifecycle Provenance (Tier 1 only)

What it records: Every state transition as an immutable onchain event.

```
fill → maturation → regauge → maturation → regauge → bottling_ready → bottled
```

Each transition logged with: cask ID, from state, to state, timestamp, metadata (gauge data if regauge, yield if bottled, etc.)

Data flow:



Why this matters: - Immutable chain of custody from fill to bottle - Solves the "receipts problem" — prove your whisky is what you say it is - Every regauge result is permanently recorded (can see volume trajectory over time) - Regulatory compliance: auditable trail of every event

Frequency: Event-driven (webhook from warehouse) + daily cron fallback to catch missed events.

Component 4: Privacy Layer — THE CORE (Feb 12+)

This is not an add-on. This is what makes the entire architecture viable for real participants.

No warehouse operator or fund will expose barrel counts, acquisition costs, supplier relationships, or portfolio composition on a public blockchain. Transparency and commercial reality are directly opposed. This is why every existing tokenized whisky project is marketing over substance — they can't solve this tension.

Confidential HTTP resolves it: Prove "cask count \geq token supply" WITHOUT revealing: - Exact number of casks - Which warehouses - Barrel acquisition costs - Supplier relationships - Portfolio composition

What doesn't need privacy: - Lifecycle events (provenance = transparency by design) - Physical attributes of YOUR casks (token holders should see what they own) - Age, type, state (not commercially sensitive)

Generalizability: This pattern — privacy-preserving proof of reserve for physically-held assets with commercially sensitive custody data — applies to any bonded warehouse, any custodian, any asset class: wine collections, art storage, precious metals vaults, commodity warehouses. Whisky casks are the demo. The infrastructure is general.

Track strategy: Primary target is Privacy track (\$16k). Less crowded field, we're early adopters of a capability that drops Feb 12, and the pattern generalizes. DeFi & Tokenization (\$20k) is the backup.

4. THE DATA SOURCE PROBLEM (HONEST)

The entire system depends on ONE data source: the warehouse API.

Current state (hackathon): - We build a bun/Hono server that serves seeded mock data - CRE workflows fetch from it - This demonstrates the architecture

What makes it credible anyway: - Warehouse operators are TTB-regulated (US) / HMRC-WOWGR regulated (UK) - Falsifying records is a federal offense - The warehouse's business model depends on reputation (custody is their product) - This isn't random API data — it's legally mandated records

Future evolution (beyond hackathon — mention in demo, don't build): - Multiple independent data sources (warehouse A + warehouse B + independent auditor) - CRE can

fetch from all three and require 2-of-3 agreement - IoT sensors (weight, temperature, humidity) as additional verification - Third-party regauge services as independent confirmation - This is an architecture that SCALES to more trust sources. Hackathon = 1 source. Production = N sources.

5. STAKEHOLDER VALUE MAP

Cask Owner / Token Issuer (Tim)

| Problem Today | What We Solve |
|---|---|
| Investors can't verify casks exist | Continuous proof of reserve, independently verified |
| Quarterly PDF appraisals are slow and expensive | Real-time physical attributes onchain |
| No provenance trail beyond paper records | Immutable lifecycle events from fill to bottle |
| Sharing inventory data exposes commercial info | Confidential HTTP: prove reserves without revealing details |
| High cost of capital due to opacity | Transparent verified data → investors accept lower risk premium |

Fund Manager / Institutional Investor

| Problem Today | What We Solve |
|---|---|
| Stale NAV between annual appraisals | Continuous attribute data — build your own model on verified inputs |
| Expensive custody audits | Proof of reserve replaces point-in-time audits |
| Chain-of-custody documentation for compliance | Lifecycle events satisfy audit trail requirements |
| Can't independently verify broker claims | Read the contract. It's all there. |

DeFi Protocol (lending, DEX, etc.)

| Problem Today | What We Solve |
|--|---|
| Can't accept RWA collateral without trusted price feed | Physical attribute oracle provides inputs for collateral models |
| No proof that collateral exists | Proof of reserve is composable — any protocol can read it |
| No standard for RWA data onchain | Standardized struct for cask attributes that any protocol can consume |

Warehouse Operator

| Problem Today | What We Solve |
|--|--|
| Manual reporting to multiple clients/investors | Report once to API, oracle distributes |
| Liability for valuation claims | We don't publish valuations. Just facts. Warehouse not opining on price. |
| Competitive risk from sharing data | Confidential HTTP: data verified without public exposure |

6. SMART CONTRACT REVISION

Based on this grounded mapping, the contract should store:

Per-Cask Attributes (replaces "NAVData")

```
struct CaskAttributes {
    CaskType caskType;           // bourbon_barrel, sherry_butt, etc.
    uint256 fillDate;           // unix timestamp
    uint256 entryGaugeVolume;   // original volume in mL (integer, no decimals)
    uint256 entryGaugeProof;    // proof * 10 (e.g., 1200 = 120 proof)
    uint256 lastRegaugeRLA;     // regauged litres of alcohol, mL
    uint256 lastRegaugeABV;     // ABV * 100 (e.g., 5850 = 58.5%)
    uint256 lastRegaugeDate;    // unix timestamp (shows data freshness)
    uint256 estimatedVolume;   // model estimate, clearly separate from gauge
    CaskState state;
    string warehouseId;
}
```

Reserve Attestation (two explicit modes)

```

struct ReserveAttestationPublic {
    uint256 physicalCaskCount;
    uint256 totalTokenSupply;
    uint256 tokensPerCask;
    uint256 reserveRatio;          // scaled 1e18, (physicalCaskCount * tokensPerCask) /
    uint256 timestamp;
    bytes32 attestationHash;
}

struct ReserveAttestationPrivate {
    bool isFullyReserved;
    uint256 timestamp;
    bytes32 attestationHash;
}

```

Privacy-track submissions use `ReserveAttestationPrivate` semantics.

Lifecycle Event (stays the same — event log)

```

event LifecycleTransition(
    uint256 indexed caskId,
    CaskState fromState,
    CaskState toState,
    uint256 timestamp,
    uint256 regaugeRLA,           // 0 if not a regauge event
    uint256 regaugeABV           // 0 if not a regauge event
)

```

Optional: Reference Valuation (clearly labeled)

```

struct ReferenceValuation {
    uint256 caskId;
    uint256 estimatedValueUsd;   // 1e18 scaled
    string methodology;         // "age_curve_v1" | "market_comp_v1" | etc.
    uint256 timestamp;
    // Anyone can ignore this. The raw attributes are right there.
}

```

7. API REVISION

Endpoints reflecting actual data tiers:

| Endpoint | Tier | Returns |
|------------------------------------|------|---|
| GET /inventory | 1 | Cask count, IDs, existence attestation |
| GET /cask/{id}/attributes | 1 | Type, fill date, entry gauge, last regauge, state, warehouse |
| GET /cask/{id}/estimate | 2 | Current volume estimate, bottle yield estimate, model version |
| GET /cask/{id}/lifecycle | 1 | State history with timestamps and gauge data |
| GET /portfolio/summary | 1+2 | Aggregate counts, volumes, age distribution |
| GET /market-data | 3 | Auction comps, index values (reference only) |
| GET /cask/{id}/reference-valuation | 3 | Model-based valuation, clearly labeled as estimate |

8. CRE WORKFLOW REVISION

Architecture has 3 components, implemented as 4 workflows to remove trigger ambiguity.

Workflow 1: Proof of Reserve

- Cron hourly
- HTTP -> /inventory (1 call)
- EVM READ -> totalMinted() (1 call)
- EVM WRITE -> reserve attestation (public mode for baseline simulation, private mode for confidential submission)
- **Budget: 1 HTTP, 1 EVM read, 1 EVM write**

Workflow 2: Physical Attribute Oracle

- Cron daily (attributes do not change fast)
- HTTP -> /portfolio/summary + /cask/{id}/attributes for recently changed casks (2 calls)
- EVM WRITE -> batch update cask attributes
- **Budget: 2 HTTP, 0 EVM reads, 1 EVM write**

Workflow 3: Lifecycle Webhook

- HTTP trigger (webhook) for real-time events
- EVM WRITE -> lifecycle event
- **Budget: 0-1 HTTP, 0 EVM reads, 1 EVM write**

Workflow 4: Lifecycle Reconcile

- Cron daily fallback to catch missed events
- HTTP -> `/cask/{id}/lifecycle` (1 call)
- EVM WRITE -> lifecycle event replay for missed transitions
- **Budget: 1 HTTP, 0 EVM reads, 1 EVM write**

All four workflows are within CRE execution limits for hackathon scope.

9. DEMO NARRATIVE (REVISED)

The pitch leads with the tension that privacy resolves:

"The whisky cask investment market is \$75B with no verifiable custody. Warehouse data is commercially sensitive — barrel counts, acquisition costs, supplier relationships. No custodian will put that on a public blockchain. So every tokenized whisky project is stuck: you can have transparency or you can have participation, but not both.

Confidential HTTP breaks that trade-off. We built CRE infrastructure that lets bonded warehouses prove reserves, attest to physical attributes, and record lifecycle events — without exposing a single piece of commercially sensitive data. The chain sees 'reserves ≥ tokens: true.' Nobody sees the book.

We built it for whisky casks because that's our business. But the pattern works for any physically-held asset with sensitive custody data — wine, art, precious metals, commodities."

10. WHAT WE BUILD FEB 6 (EXECUTION ORDER)

1. **Smart contract** (day 1-2): CaskAttributes struct, ReserveAttestationPublic/Private semantics, LifecycleTransition events, onReport routing, read functions
2. **Mock warehouse API** (day 2-3): Bun/Hono serving seeded cask data across all endpoints, clearly separated tiers

3. **CRE workflows** (day 3-5): PoR, attributes oracle, lifecycle webhook + lifecycle reconcile.
Get simulation working end-to-end.
4. **Privacy layer** (Feb 12+): Confidential HTTP swap on PoR inventory fetch
5. **Reference valuation model** (if time): Optional Tier 3 layer, clearly labeled and offchain-consumable
6. **Demo video** (Feb 25+): Record terminal demos, architecture diagram, narration